

Intro to: Inter-Processor Communications (IPC)

Agenda

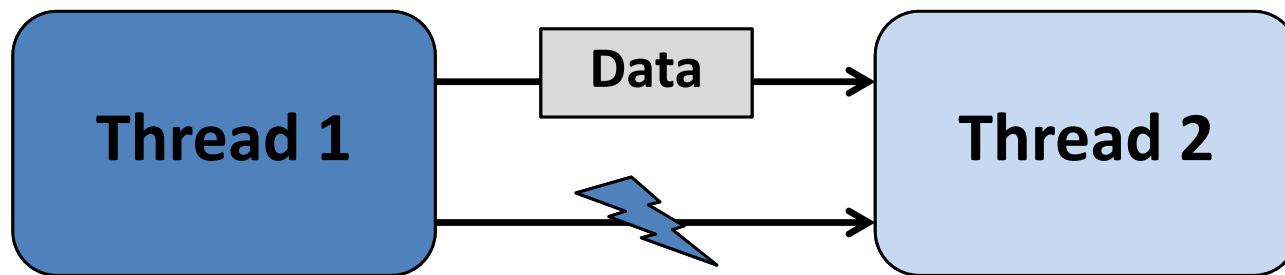
- ◆ Basic Concepts
- ◆ IPC Services
- ◆ Setup and Examples
- ◆ IPC Transports
- ◆ For More Information

Basic Concepts

- ◆ **Basic Concepts**
- ◆ IPC Services
- ◆ Setup and Examples
- ◆ IPC Transports
- ◆ For More Information

IPC – Definition

- ◆ IPC = Inter-Processor Communication
- ◆ While this definition is rather generic, it really means:
“Transporting data and/or signals between threads of execution”
- ◆ These threads could be located anywhere:



CorePac 0 —————> CorePac 0

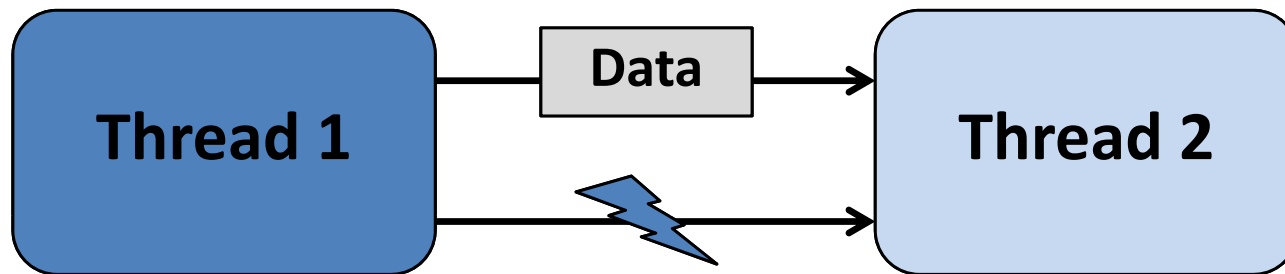
CorePac 0 —————> CorePac 1

Device 0 —————> Device 1

How would YOU solve this problem?

IPC – Possible Solutions

- ◆ How do you transport the data and signal?
 - **Manual:** Develop your own protocol using device resources
 - **Auto:** Using existing RTOS/Framework Utilities (i.e., IPC)



CorePac 0 —————> CorePac 0

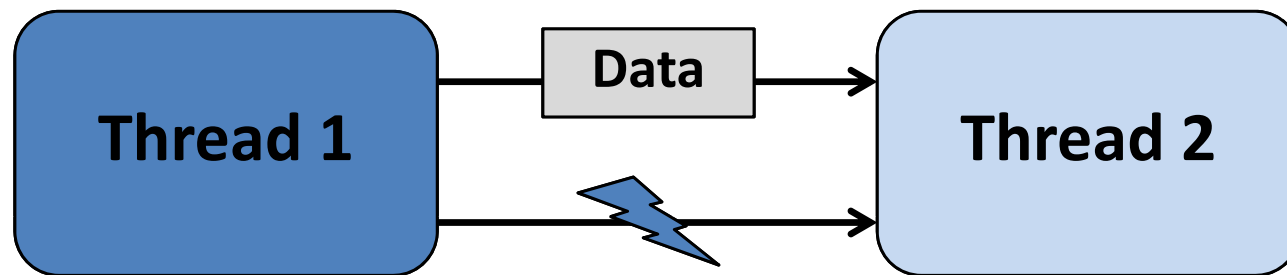
CorePac 0 —————> CorePac 1

Device 0 —————> Device 1

What solutions exist in TI's RTOS to perform IPC ?

IPC – RTOS/Framework Solutions

- ◆ **SAME CorePac:** TI's RTOS (SYS/BIOS) supports several services for inter-thread communication (*e.g. semaphores, queues, mailboxes, etc.*).
- ◆ **DIFFERENT CorePac:** The IPC framework supports communications between CorePacs via several transports.
- ◆ **DIFFERENT DEVICE:** IPC transports can also be implemented between devices.
- ◆ **KEY:** Same IPC APIs can be used for local or remote communications.

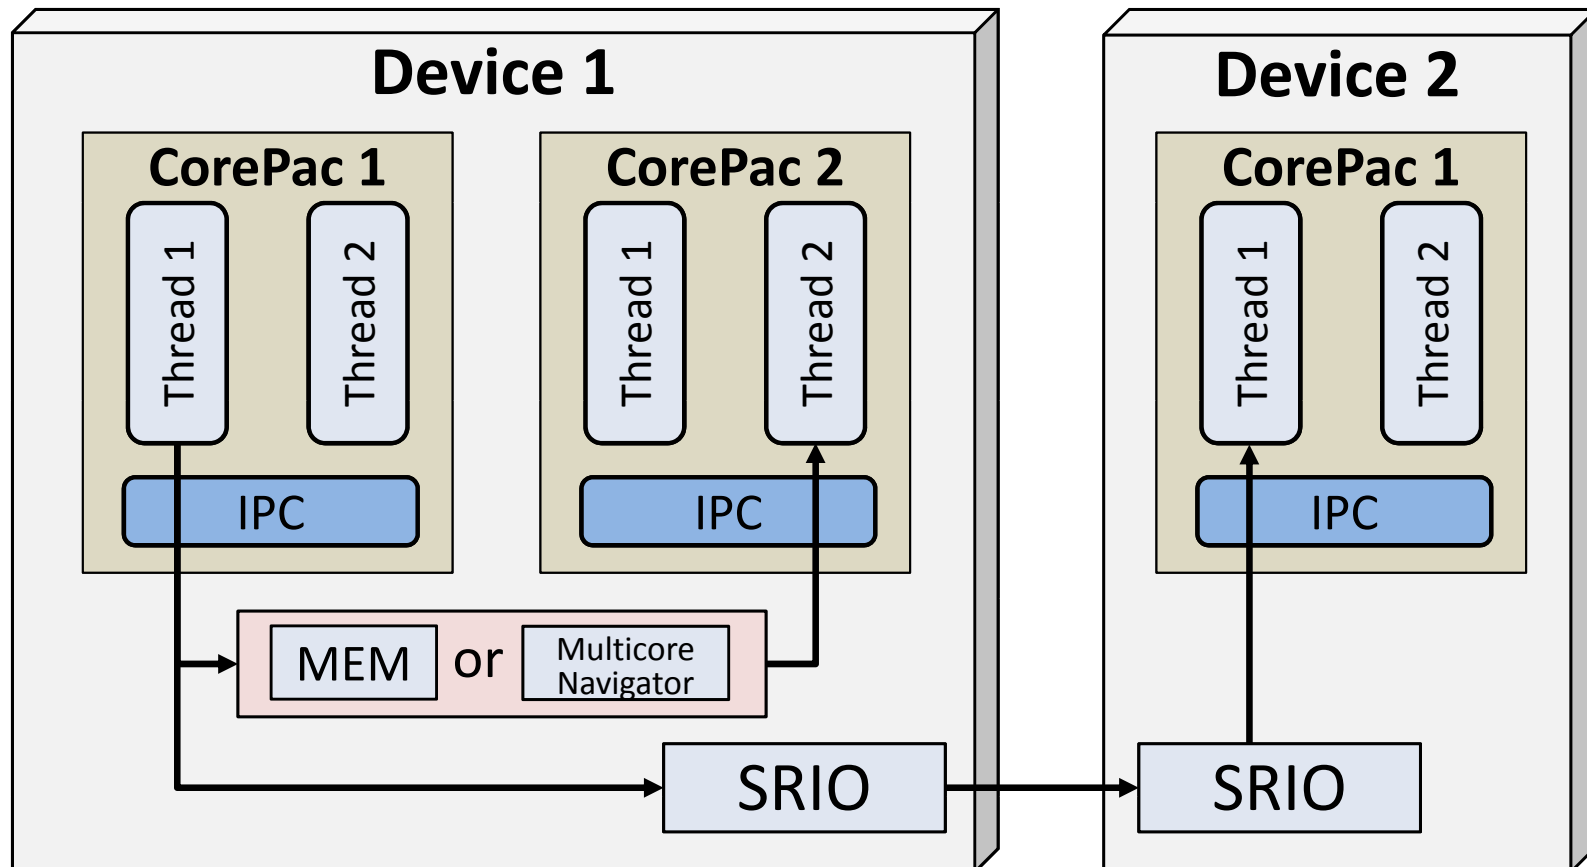


Solutions

CorePac 0	—————>	CorePac 0	➤ SYS/BIOS (or IPC)
CorePac 0	—————>	CorePac 1	➤ IPC + transport
Device 0	—————>	Device 1	➤ IPC + transport

IPC – Transports

- ◆ Current IPC implementation uses several transports:
 - CorePac → CorePac (Shared Memory Model, Multicore Navigator)
 - Device → Device (Serial Rapid I/O)
- ◆ Chosen at configuration; Same code regardless of thread location.

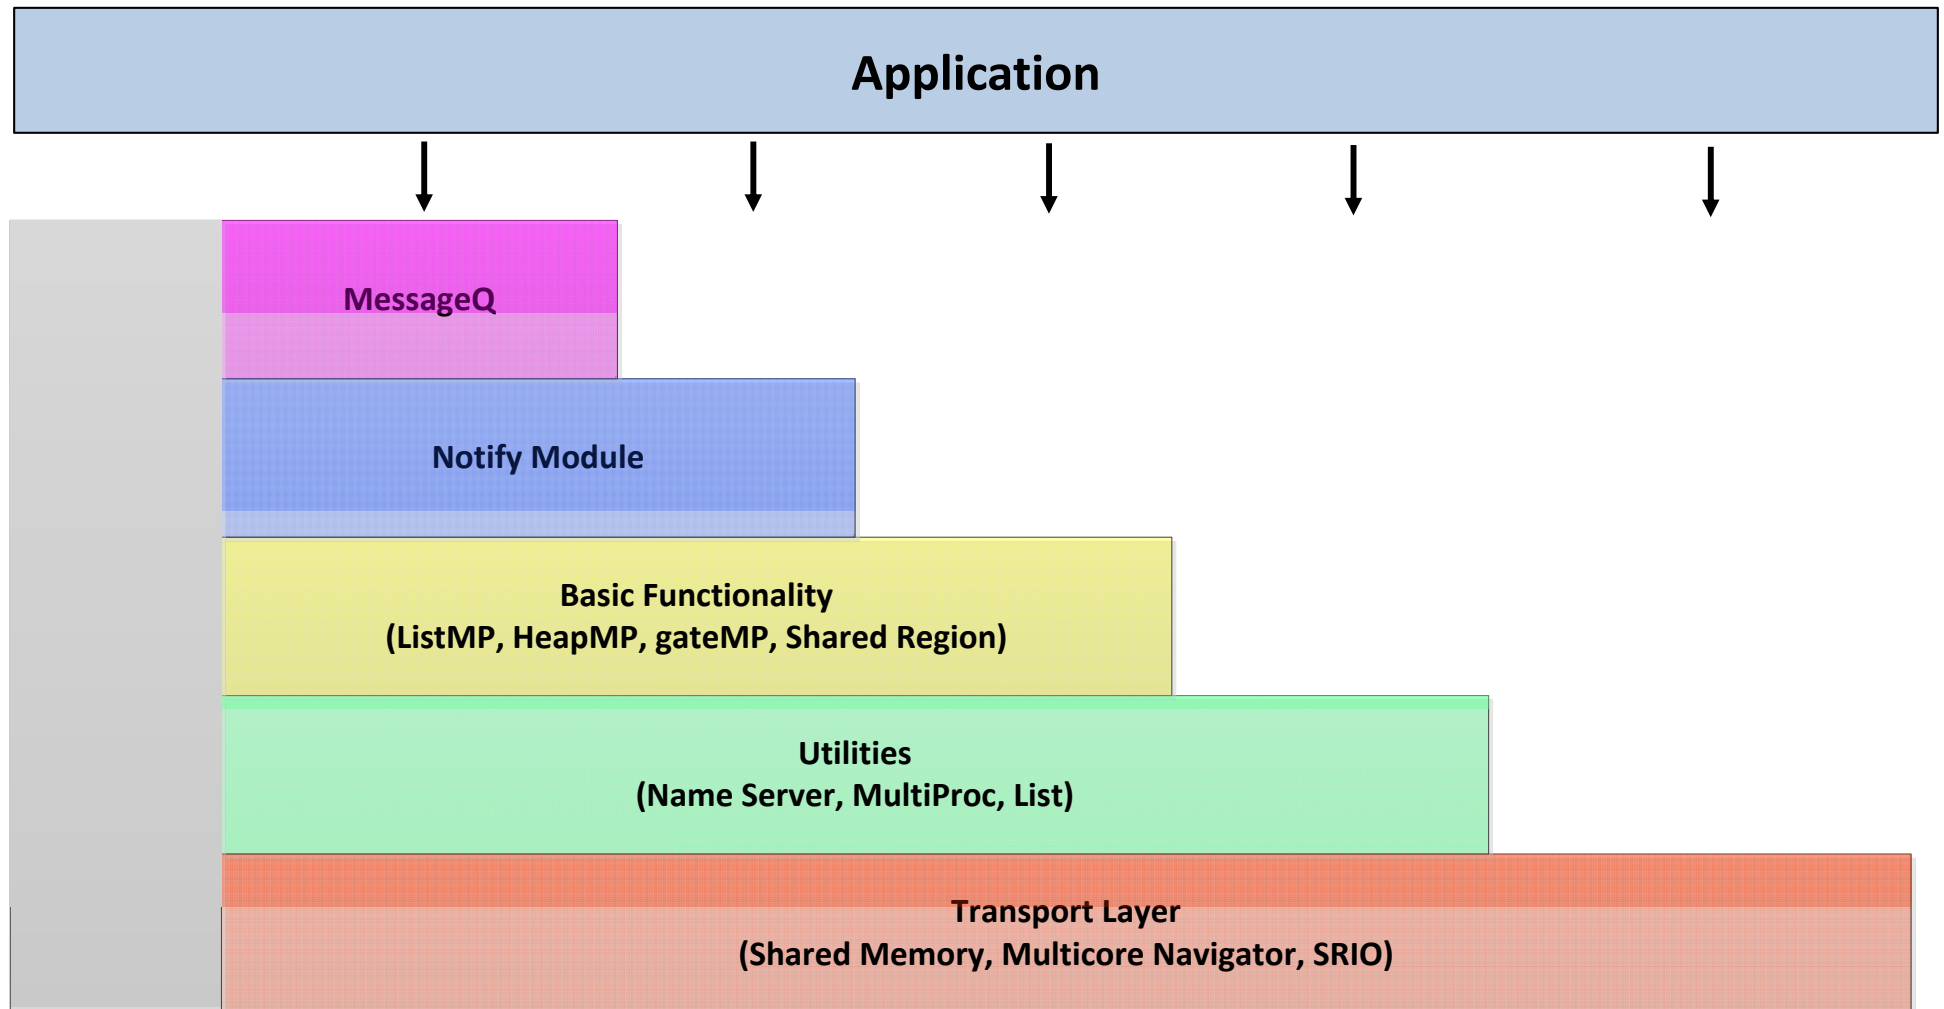


IPC Services

- ◆ Basic Concepts
- ◆ **IPC Services**
 - Message Queue
 - Notify
 - Data Passing
 - Support Utilities
- ◆ Setup and Examples
- ◆ IPC Transports
- ◆ For More Information

IPC Services

- ◆ The IPC package is a set of APIs
- ◆ MessageQ uses the modules below ...
- ◆ But each module can also be used independently.



IPC Services – Message Queue

- ◆ Basic Concepts
- ◆ **IPC Services**
 - **Message Queue**
 - Notify
 - Data Passing
 - Support Utilities
- ◆ Setup and Examples
- ◆ IPC Transports
- ◆ For More Information

MessageQ – Highest Layer API

- ◆ SINGLE reader, multiple WRITERS model (READER owns queue/mailbox)
- ◆ Supports structured sending/receiving of variable-length messages, which can include (pointers to) data
- ◆ Uses all of the IPC services layers along with IPC Configuration & Initialization
- ◆ APIs do not change if the message is between two threads:
 - ◆ On the same core
 - ◆ On two different cores
 - ◆ On two different devices
- ◆ APIs do NOT change based on transport – only the CFG (init) code
 - ◆ Shared memory
 - ◆ Multicore Navigator
 - ◆ SRIO

MessageQ and Messages

- Q How does the writer connect with the reader queue?
 - A MultiProc and name server keep track of queue names and core IDs.
- Q What do we mean when we refer to structured messages with variable size?
 - A Each message has a standard header and data. The header specifies the size of payload.
- Q How and where are messages allocated?
 - A List utility provides a double-link list mechanism. The actual allocation of the memory is done by HeapMP, SharedRegion, and ListMP.
- Q If there are multiple writers, how does the system prevent race conditions (e.g., two writers attempting to allocate the same memory)?
 - A GateMP provides hardware semaphore API to prevent race conditions.
- Q What facilitates the moving of a message to the receiver queue?
 - A This is done by Notify API using the transport layer.
- Q Does the application need to configure all these modules?
 - A No. Most of the configuration is done by the system. More details later.

Using MessageQ (1/3)

CorePac 2 - READER



```
MessageQ_create("myQ", *synchronizer);  
MessageQ_get("myQ", &msg, timeout);
```

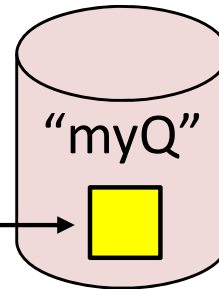
- ◆ MessageQ transactions begin with **READER** creating a MessageQ.
- ◆ **READER's** attempt to get a message results in a block (unless timeout was specified), since no messages are in the queue yet.

What happens next?

Using MessageQ (2/3)

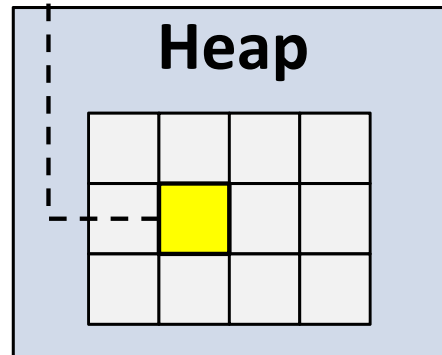
CorePac 1 - WRITER

```
MessageQ_open ("myQ", ...);  
msg = MessageQ_alloc (heap, size,...);  
MessageQ_put("myQ", msg, ...);
```



CorePac 2 - READER

```
MessageQ_create("myQ", ...);  
MessageQ_get("myQ", &msg...);
```



- ◆ **WRITER** begins by opening MessageQ created by **READER**.
- ◆ **WRITER** gets a message block from a heap and fills it, as desired.
- ◆ **WRITER** puts the message into the MessageQ.

How does the **READER** get unblocked?

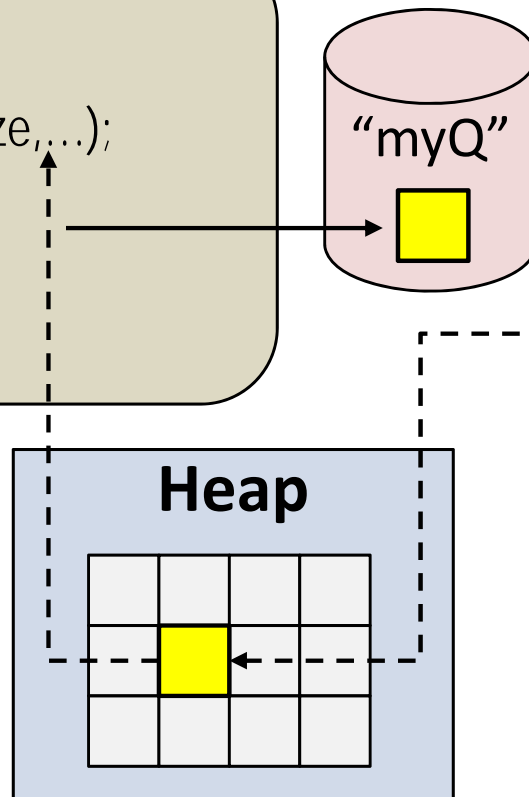
Using MessageQ (3/3)

CorePac 1 - WRITER

```
MessageQ_open ("myQ", ...);  
msg = MessageQ_alloc (heap, size,...);  
MessageQ_put("myQ", msg, ...);  
MessageQ_close("myQ", ...);
```

CorePac 2 - READER

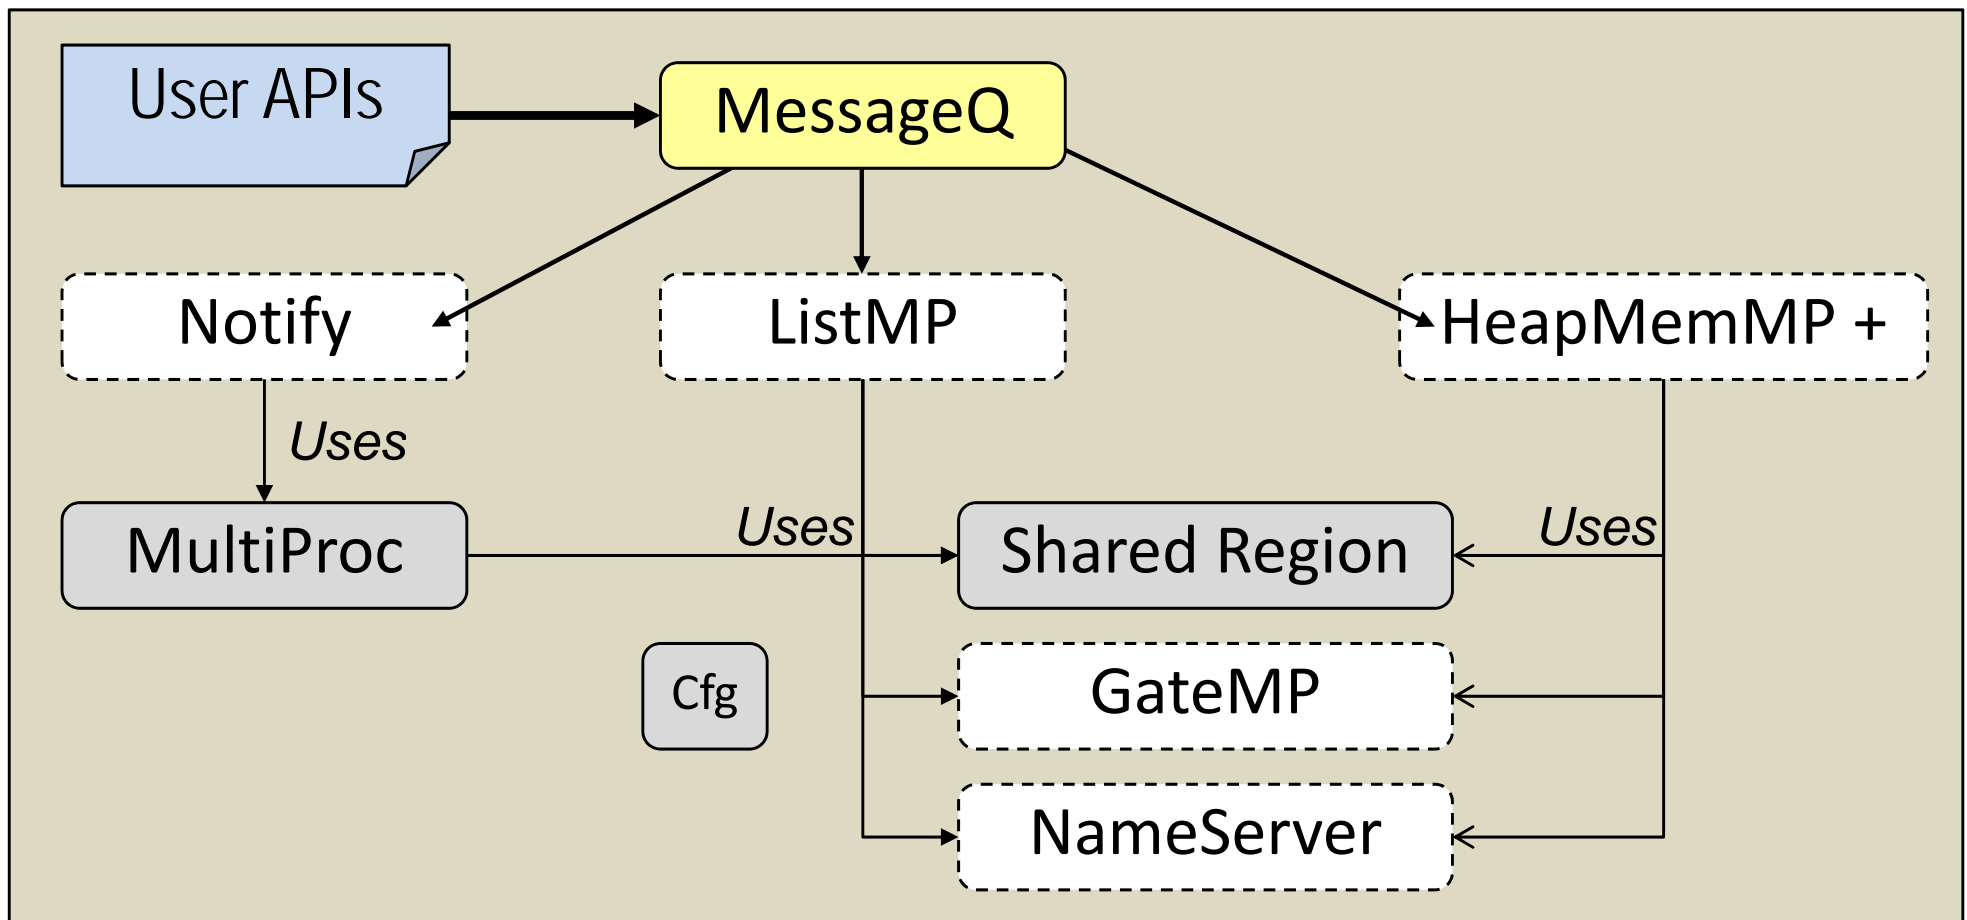
```
MessageQ_create("myQ", ...);  
MessageQ_get("myQ", &msg...);  
*** PROCESS MSG ***  
MessageQ_free("myQ", ...);  
MessageQ_delete("myQ", ...);
```



- ◆ Once **WRITER** puts msg in MessageQ, **READER** is unblocked.
- ◆ **READER** can now read/process the received message.
- ◆ **READER** frees message back to Heap.
- ◆ **READER** can optionally delete the created MessageQ, if desired.

MessageQ – Configuration

- ◆ All API calls use the MessageQ module in IPC.
- ◆ User must also configure MultiProc and SharedRegion modules.
- ◆ All other configuration/setup is performed automatically by MessageQ.



MessageQ – Miscellaneous Notes

- ◆ O/S independent:
 - ◆ If one CorePac is running LINUX and using SysLink, the API calls do not change.
 - ◆ SysLink is runtime software that provides connectivity between processors (running Linux, SYSBIOS, etc.)
- ◆ Messages can be allocated statically or dynamically.
- ◆ Timeouts are allowed when a task receives a message.
- ◆ User can specify three priority levels:
 - ◆ Normal
 - ◆ High
 - ◆ Urgent

More Information About MessageQ

- ◆ All structures and function descriptions can be found within the release:

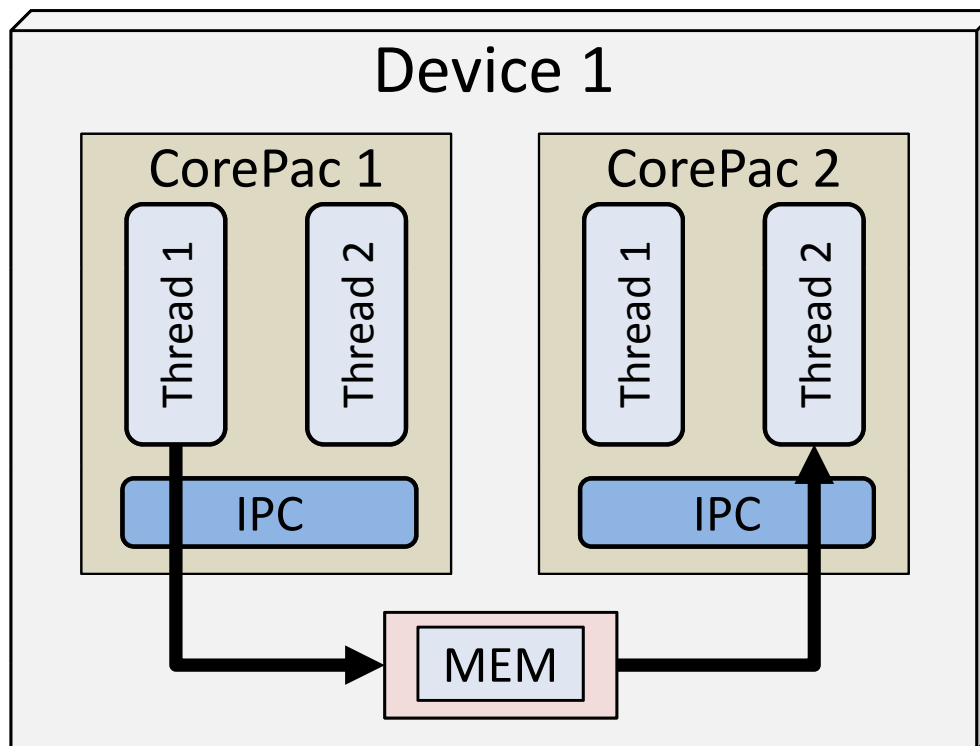
`\ipc_U_ZZ_YY_XX\docs\doxygen\html_message_q_8h.html`

IPC Services - Notify

- ◆ Basic Concepts
- ◆ **IPC Services**
 - Message Queue
 - **Notify**
 - Data Passing
 - Support Utilities
- ◆ Setup and Examples
- ◆ IPC Transports
- ◆ For More Information

Using Notify – Concepts

- ◆ In addition to moving MessageQ messages, Notify:
 - ◆ Can be used independently of MessageQ
 - ◆ Is a simpler form of IPC communication
 - ◆ Is used ONLY with shared memory transport



Notify Model

- ◆ Comprised of **SENDER** and **RECEIVER**.
- ◆ The **SENDER** API requires the following information:
 - ◆ Destination (**SENDER** ID is implicit)
 - ◆ 16-bit Line ID
 - ◆ 32-bit Event ID
 - ◆ 32-bit payload (For example, a pointer to message handle)
- ◆ The **SENDER** API generates an interrupt (an event) in the destination.
- ◆ Based on Line ID and Event ID, the **RECEIVER** schedules a pre-defined call-back function.

Notify Model

Sender

Receiver

During initialization, link the Event ID and Line ID with call back function.

```
Void Notify_registerEvent(srcId, lineId, eventId, cbFxn, cbArg);
```

During run time, the Sender sends a notify to the Receiver.

```
Void Notify_sendEvent(dstId, lineId, eventId, payload, waitClear);
```

INTERRUPT

Based on the Sender's Source ID, Event ID, and Line ID,
the call-back function that was registered during initialization is called with the argument payload.

Notify Implementation

- Q How are interrupts generated for shared memory transport?
 - A The IPC hardware registers are a set of 32-bit registers that generate interrupts. There is one register for each core.
- Q How are the notify parameters stored?
 - A List utility provides a double-link list mechanism. The actual allocation of the memory is done by HeapMP, SharedRegion, and ListMP.
- Q How does the notify know to send the message to the correct destination?
 - A MultiProc and name server keep track of the core ID.
- Q Does the application need to configure all these modules?
 - A No. Most of the configuration is done by the system.

Example Callback Function

```
/*  
 * ===== cbFxn =====  
 * This fxn was registered with Notify. It is called when any event is sent to this CPU.  
 */  
UInt32 recvProcid ;  
UInt32 seq  ;  
void cbFxn(UInt16 procid, UInt16 lineId, UInt32 eventId, UArg arg, UInt32 payload)  
{  
    /* The payload is a sequence number. */  
    recvProcid = procid;  
    seq = payload;  
    Semaphore_post(semHandle);  
}
```


More Information About Notify

- ◆ All structures and function descriptions can be found within the release:

`\ipc_U_ZZ_YY_XX\docs\doxygen\html_notify_8h.html`

IPC Services - Data Passing

- ◆ Basic Concepts
- ◆ **IPC Services**
 - Message Queue
 - Notify
 - **Data Passing**
 - Support Utilities
- ◆ Setup and Examples
- ◆ IPC Transports
- ◆ For More Information

Data Passing Using Shared Memory (1/2)

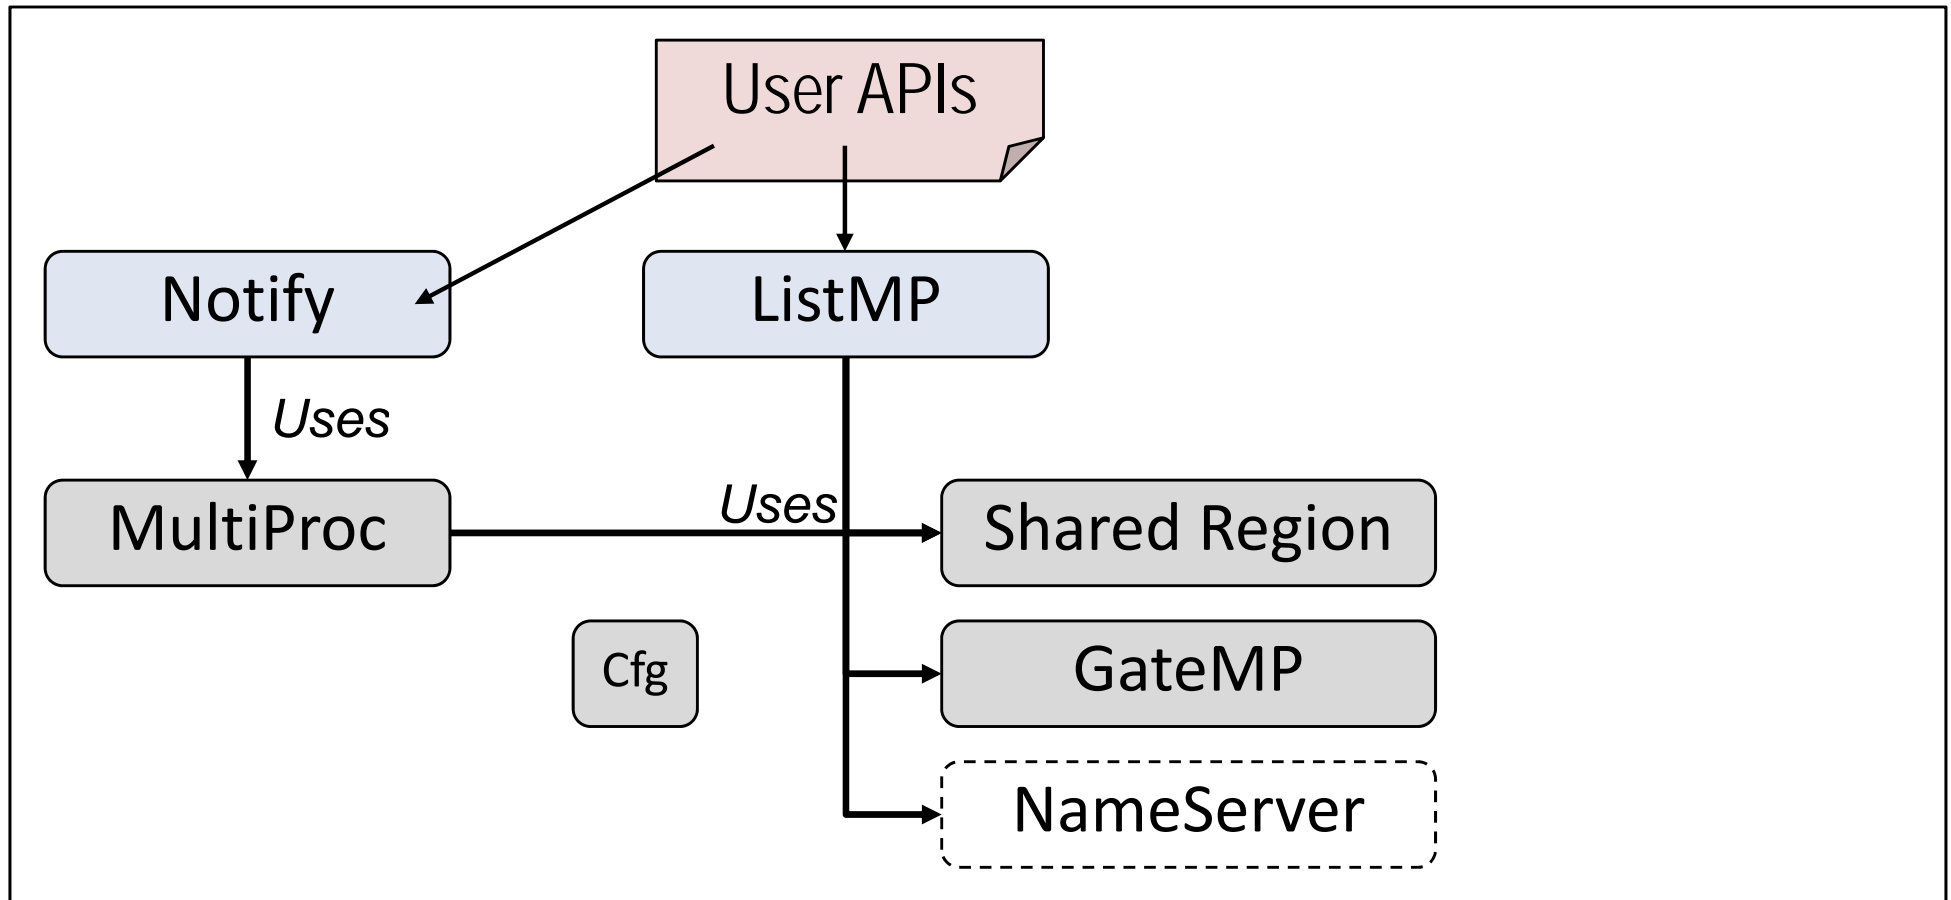
- ◆ When there is a need to allocate memory that is accessible by multiple cores, shared memory is used.
- ◆ However, the MPAX register for each core might assign a different logical address to the same physical shared memory address.
- ◆ The Shared Region module provides a translation look-up table that resolves the logical/physical address issue without user intervention.

Data Passing Using Shared Memory (2/2)

- ◆ Messages are created and freed, but not necessarily in consecutive order:
 - ◆ HeapMP provides a dynamic heap utility that supports create and free based on double link list architecture.
 - ◆ ListMP provides a double link list utility that makes it easy to create and free messages for static memory. It is used by the HeapMP for dynamic cases.
- ◆ To protect the above utilities from race conditions (e.g., multiple cores try to create messages at the same time):
 - ◆ GateMP provides hardware semaphore protection.
 - ◆ GateMP can also be used by non-IPC applications to assign hardware semaphores.

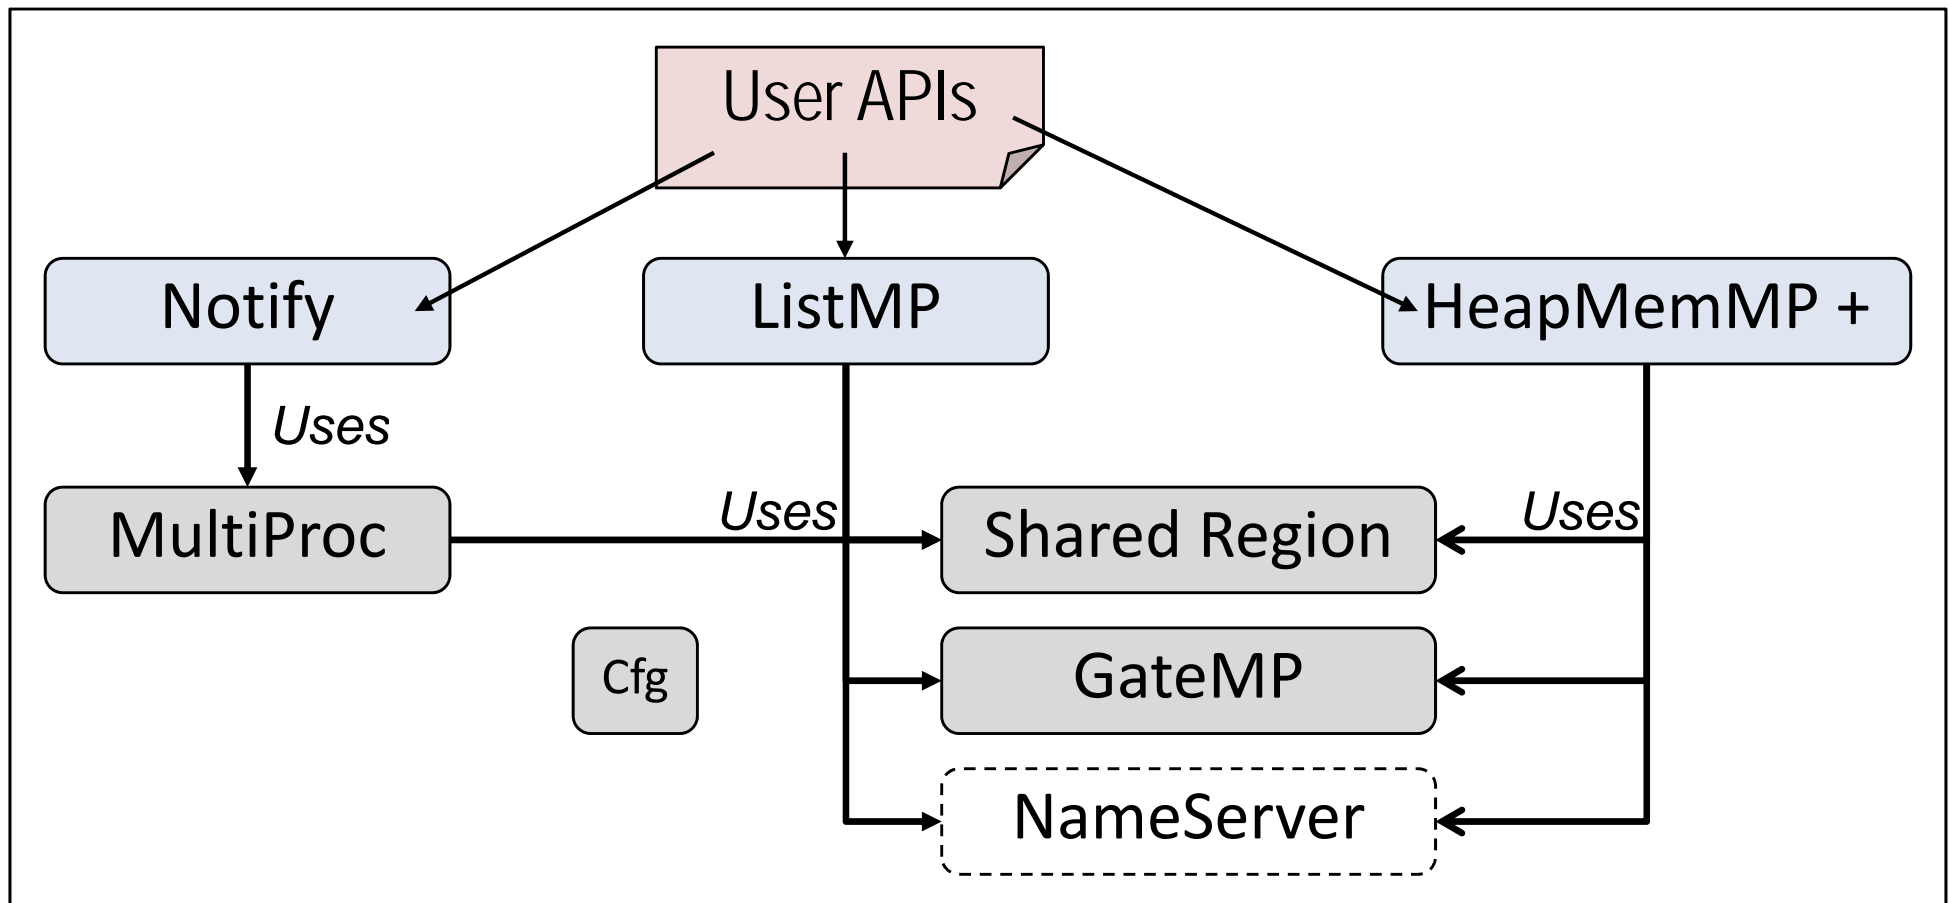
Data Passing – Static

- ◆ Data Passing uses a *double linked list* that can be shared between CorePacs; Linked list is defined **STATICALLY**.
- ◆ ListMP handles address translation and cache coherency.
- ◆ GateMP protects read/write accesses.
- ◆ ListMP is typically used by MessageQ not by itself.



Data Passing – Dynamic

- ◆ Data Passing uses a *double linked list* that can be shared between CPUs. Linked list is defined **DYNAMICALLY** (via heap).
- ◆ Same as previous, except linked lists are allocated from Heap
- ◆ Typically not used alone – but as a building block for MessageQ



Data Passing: Additional Transports

- ◆ Multicore Navigator
 - ◆ Messages are assigned to descriptors.
 - ◆ Monolithic descriptors provide a pointer to the message.
- ◆ SRIO Type 11 has a TI-developed message protocol that moves data and interrupts between devices.
- ◆ Implementation of these transports will be discussed later.

IPC Services – Support Utilities

- ◆ Basic Concepts
- ◆ **IPC Services**
 - Message Queue
 - Notify
 - Data Passing
 - **Support Utilities**
- ◆ Setup and Examples
- ◆ IPC Transports
- ◆ For More Information

IPC Support Utilities (1/2)

- ◆ IPC contains several utilities, most of which do NOT need to be configured by the user.
- ◆ Here is a short description of each IPC utility:

IPC

- Initializes IPC subsystems:
 - Applications using IPC must call `IPC_start()`.
 - In the configuration file, `setupNotify` and `setupMessageQ` specify whether to set up these IPC modules.

MultiProc

- Manages processor IDs and core names

SharedRegion

- Manages shared memory using HeapMemMP allocator
- Handles address translation for shared memory regions

[More utilities...](#)

IPC Support Utilities (2/2)

- ◆ IPC contains several utilities, most of which do NOT need to be configured by the user.
- ◆ Here is a short description of each IPC utility:

ListMP

- Provides linked list in shared memory
- Uses multi-processor gate (GateMP) to prevent collisions on lists

GateMP

- Multi-processor gate that provides local (against other threads on local core) and remote context protection

HeapMemMP

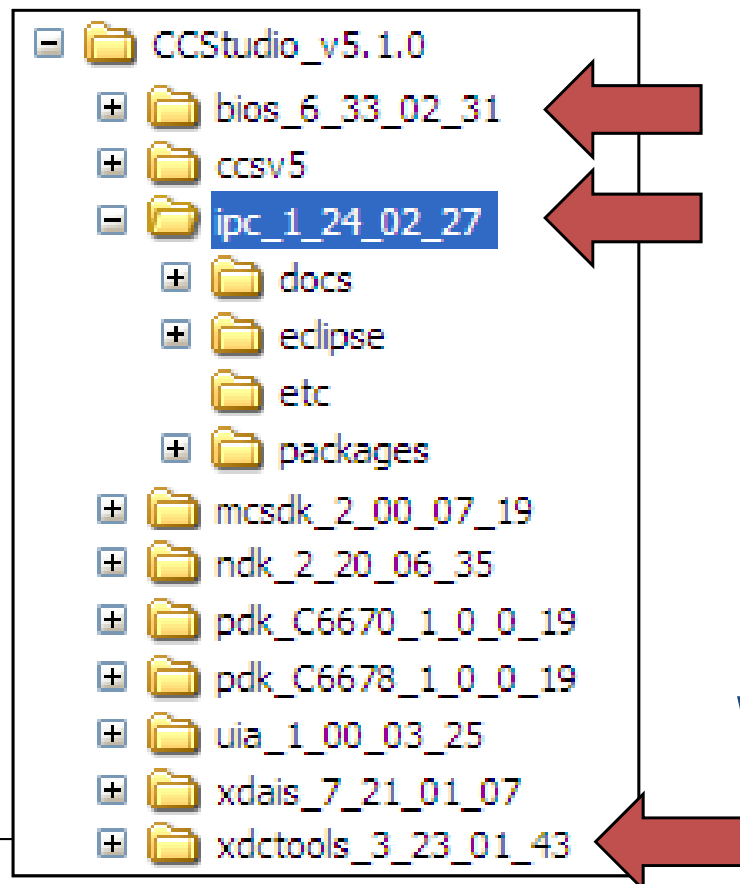
- Traditional heap; Supports variable-sized alloc/free
- All allocations are aligned on cache line boundaries.

Setup and Examples

- ◆ Basic Concepts
- ◆ IPC Services
- ◆ **Setup and Examples**
- ◆ IPC Transports
- ◆ For More Information

IPC – Tools/Setup Required

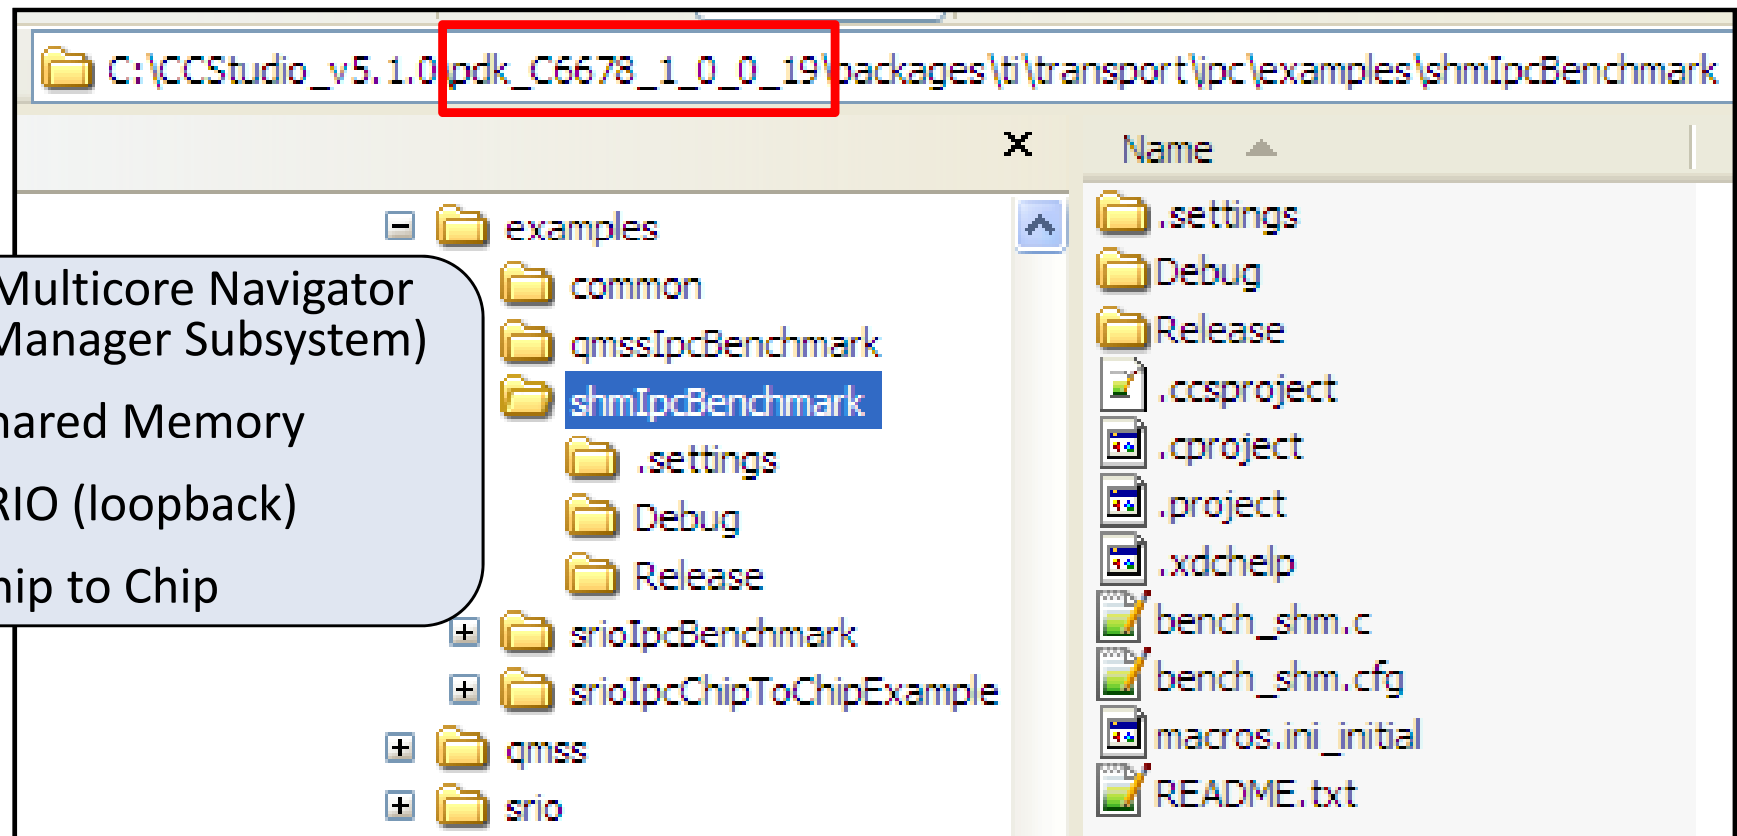
- ◆ IPC is a package (library) that is installed with the MCSDK.
Note: Installing IPC independent of MCSDK is strongly discouraged.
- ◆ IPC also requires SYS/BIOS (threading) and XDC tools (packaging) – installed with the MCSDK (supported by SYS/BIOS ROV and RTA).
- ◆ IPC is supported on the latest KeyStone multicore devices (C667x, C665x), as well as C647x and ARM+DSP devices.



What IPC examples exist in the MCSDK?

IPC – Examples

- ◆ IPC examples are installed along with the MCSDK in the PDK folder.
- ◆ Simply import these projects into CCS and analyze them.
- ◆ Some users start with this example code and modify as necessary.



IPC Transports

- ◆ Basic Concepts
- ◆ IPC Services
- ◆ Setup and Examples
- ◆ **IPC Transports**
- ◆ For More Information

IPC Transports – Intro

- ◆ An IPC **TRANSPORT** is a ...

“combination of physical H/W and *driver code* that allows two threads to communicate on the same device or across devices.”

- ◆ IPC supports three different transports:

Shared Memory

- This is the default transport.
- Uses on-chip shared memory resources and interrupt lines to signal the availability of data

Multicore Navigator

- Available on all KeyStone SoC devices
- Uses queues and descriptors plus built-in signaling to transmit/receive data and messages between threads

SRIO

- Hardware serial peripheral that connects two or more DEVICES together

- ◆ For MessageQ, only the configuration (init) code changes. All other code (e.g., put/get) remains unchanged.

Data Passing Using Multicore Navigator

The **Multicore Navigator**:

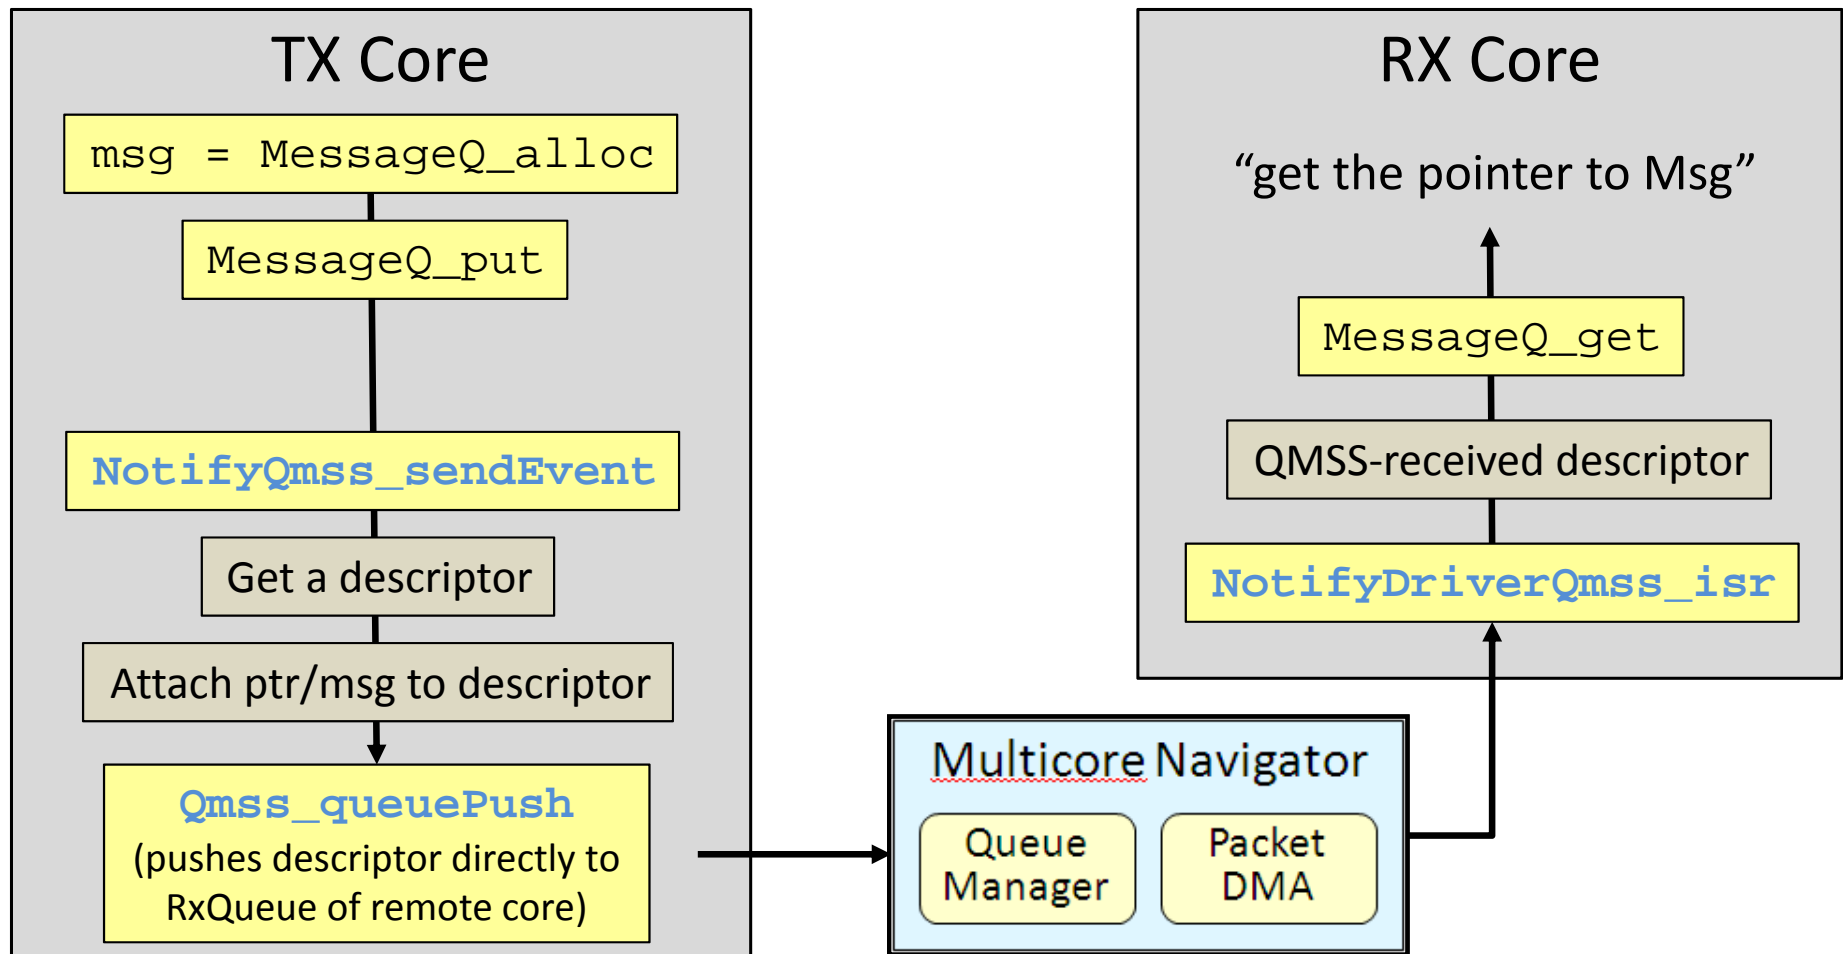
- ◆ Is an innovative packet-based infrastructure that facilitates data movement and multicore control
- ◆ Provides a highly efficient inter-core communication mechanism
- ◆ Provides H/W queues and packet DMA that are the basic building blocks for IPC

Examples of required init/CFG code are provided with the MCSDK.

Now consider the Multicore Navigator implementation ...

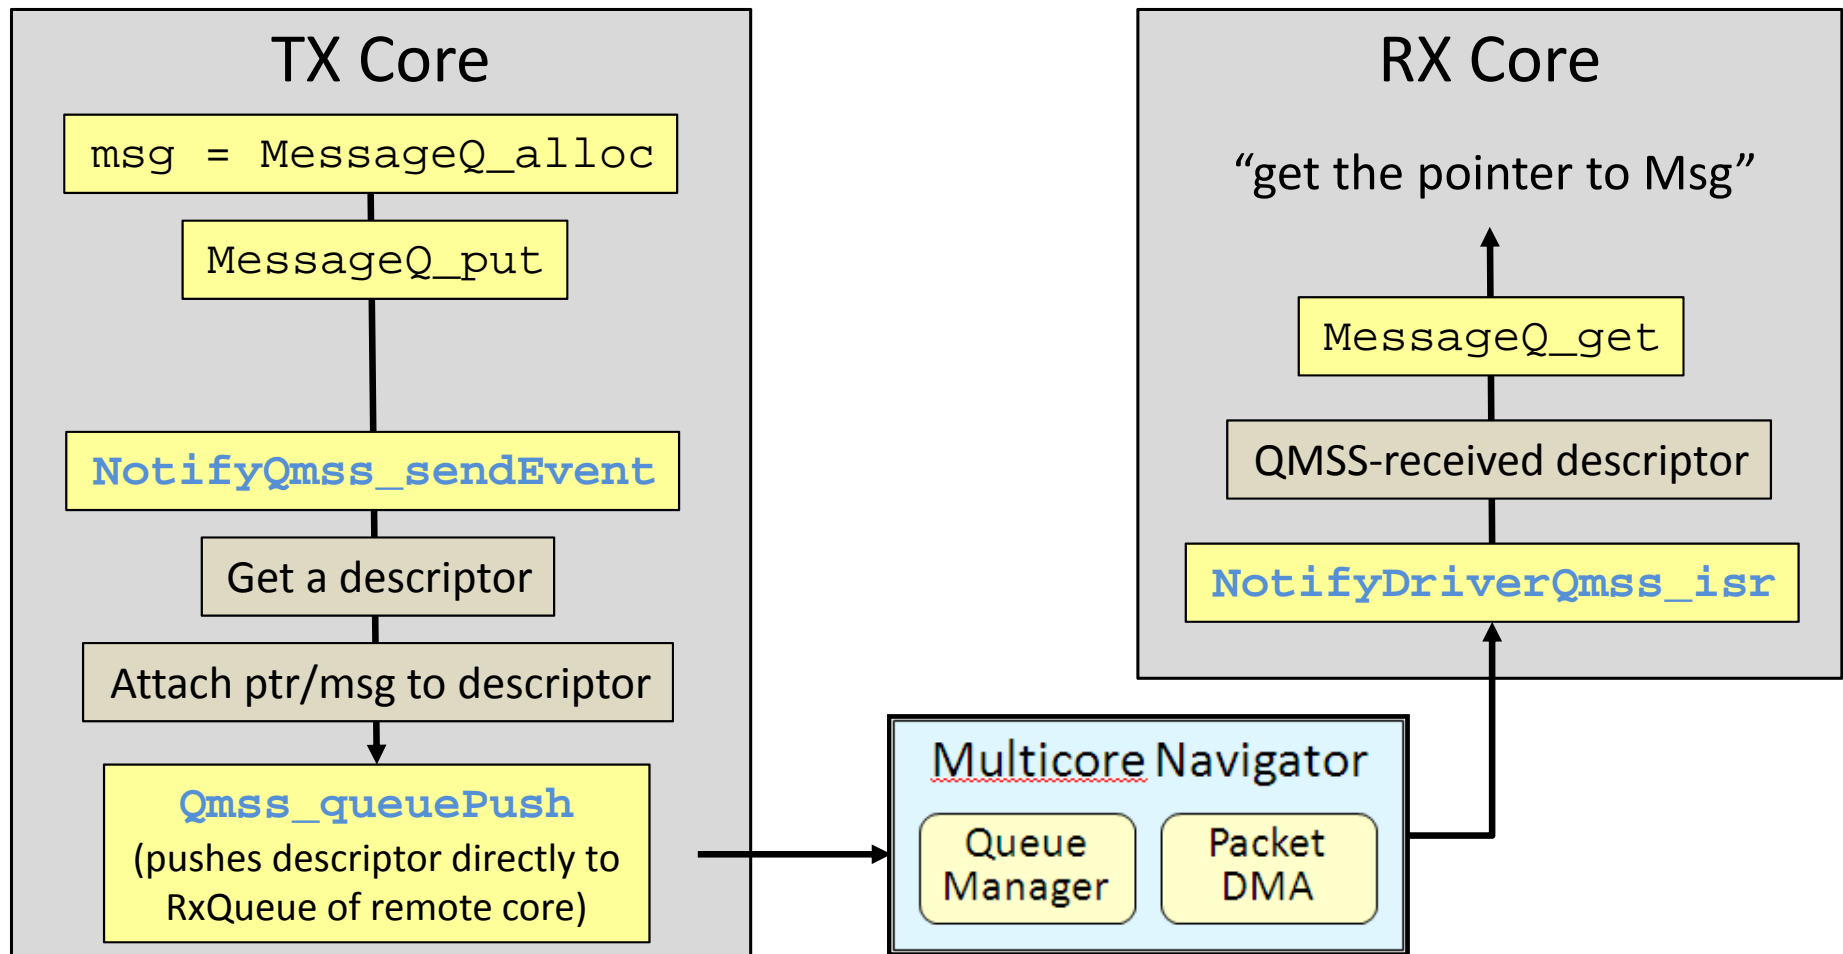
IPC Transports – Multicore Navigator (1/3)

- ◆ MessageQ_alloc allocates a heapMP buffer for the MessageQ_msg from MSMC
- ◆ The buffer contains the MessageQ_Header plus the data.
- ◆ When MessageQ_Put is called with the MessageQ_msg, it also calls the TransportQmss function **NotifyQmss_sendEvent**.



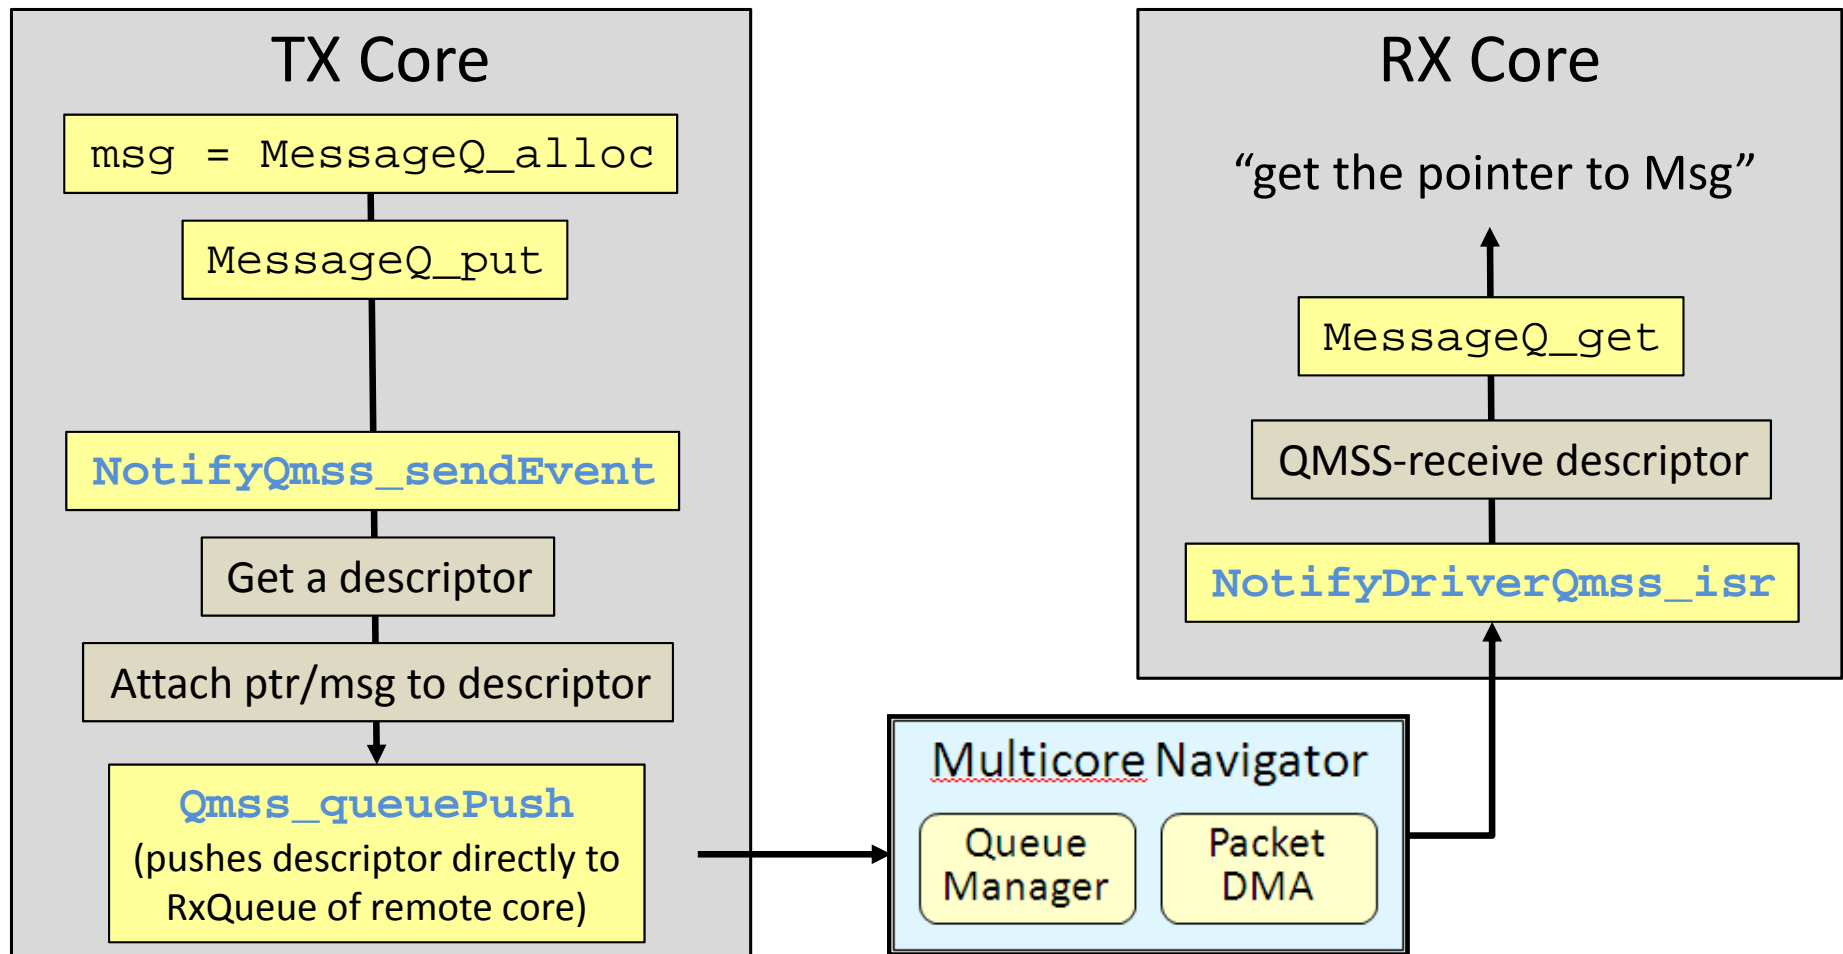
IPC Transports – Multicore Navigator (2/3)

- ◆ The TransportQmss function `NotifyQmss_sendEvent`
 - ◆ gets a descriptor
 - ◆ Adds the pointer to the `MessageQ_msg` (located in MSMC and containing the data) to the descriptor.
 - ◆ `Qmss_queuePush` pushes the descriptor into a queue.



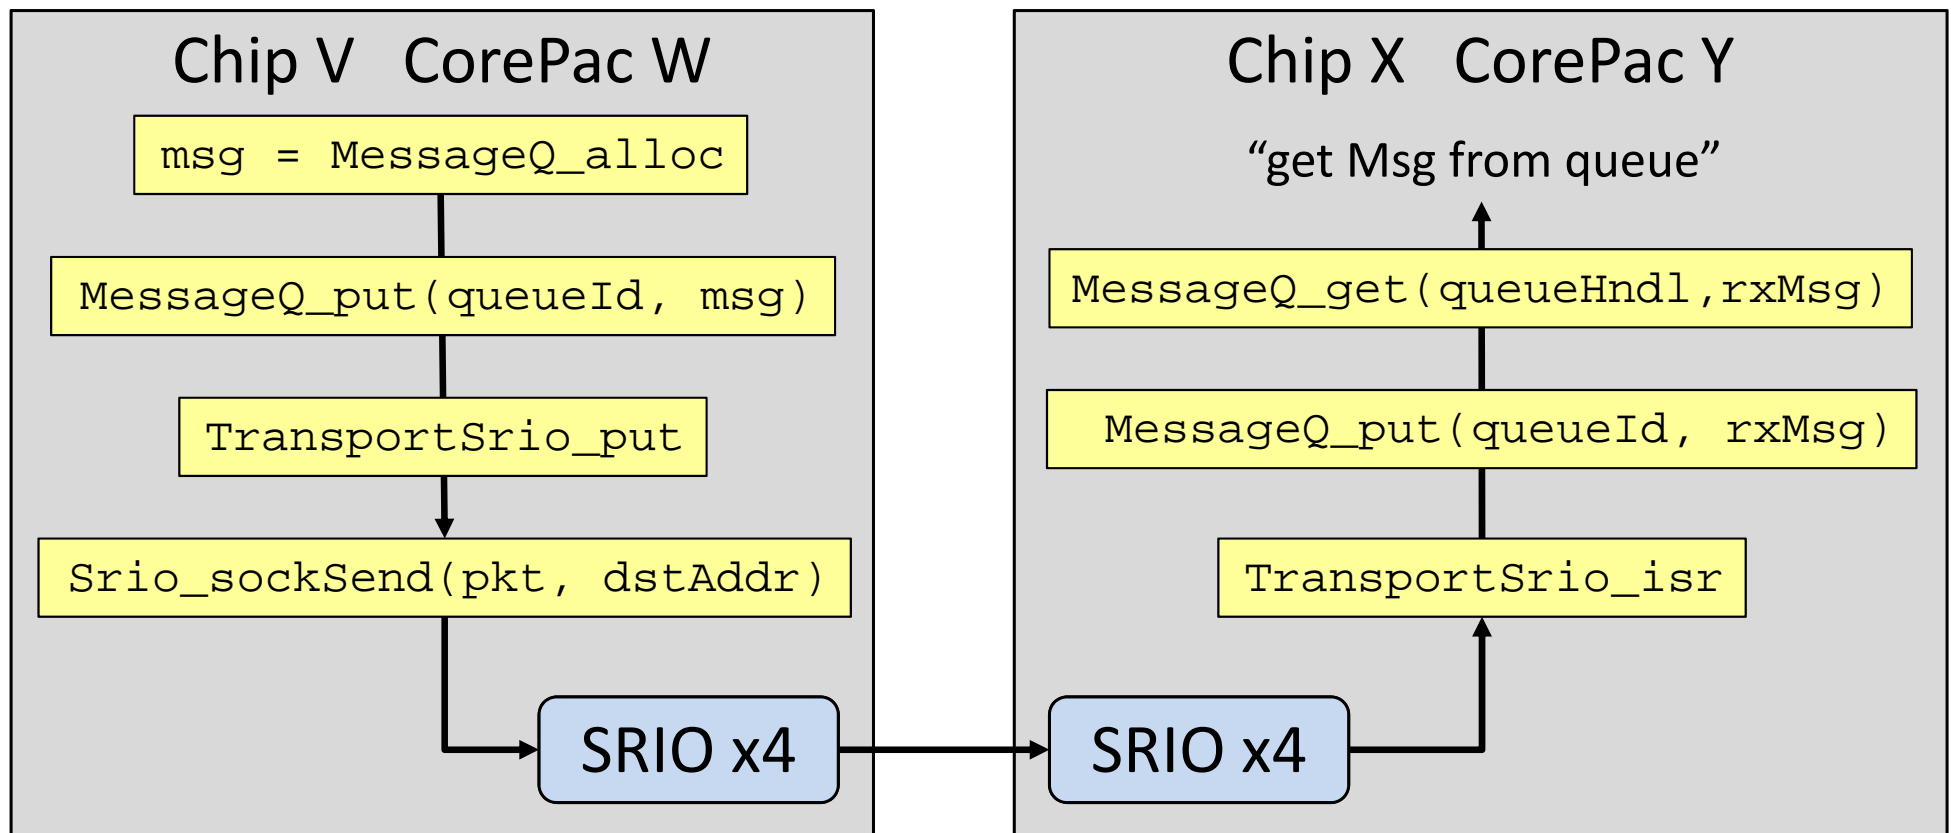
IPC Transports – Multicore Navigator (3/3)

- ◆ The Interrupt generator queue calls `notifyDriverQmss_isr`.
- ◆ The ISR schedules a QMSS receive descriptor that sends a message to the queue of Receiver.



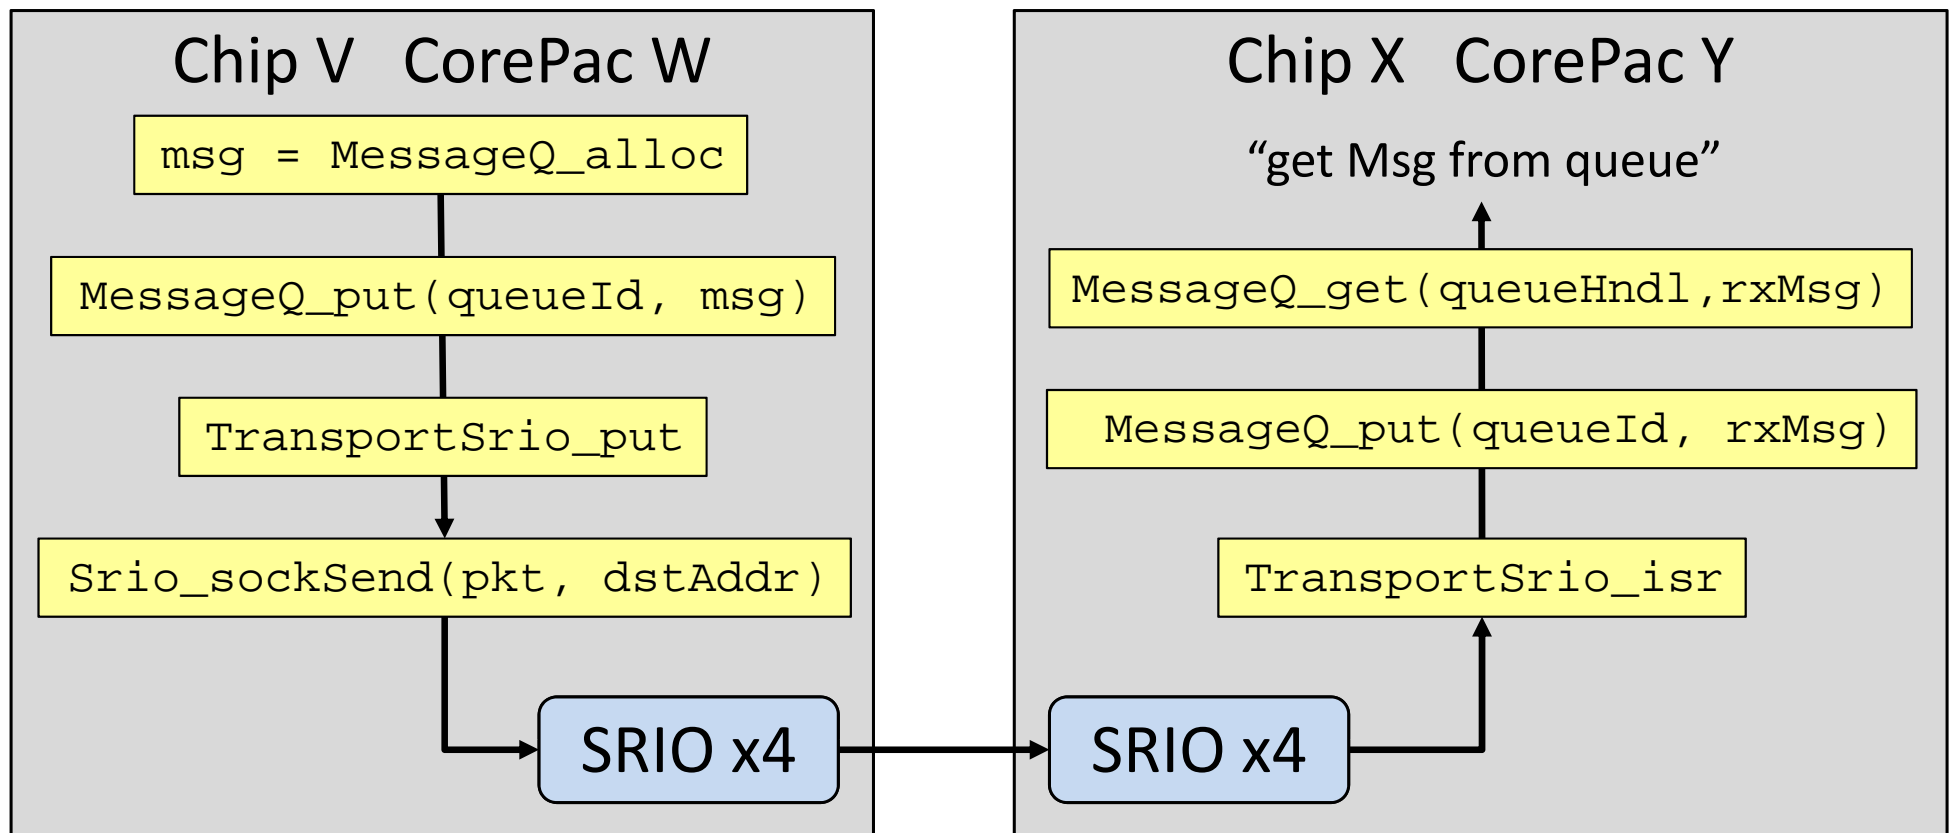
IPC Transports – SRIO (1/3)

- ◆ The **SRIO** (Type 11) transport enables MessageQ to send data between tasks, cores and devices via the SRIO IP block.
- ◆ Refer to the MCSDK examples for setup code required to use MessageQ over this transport.



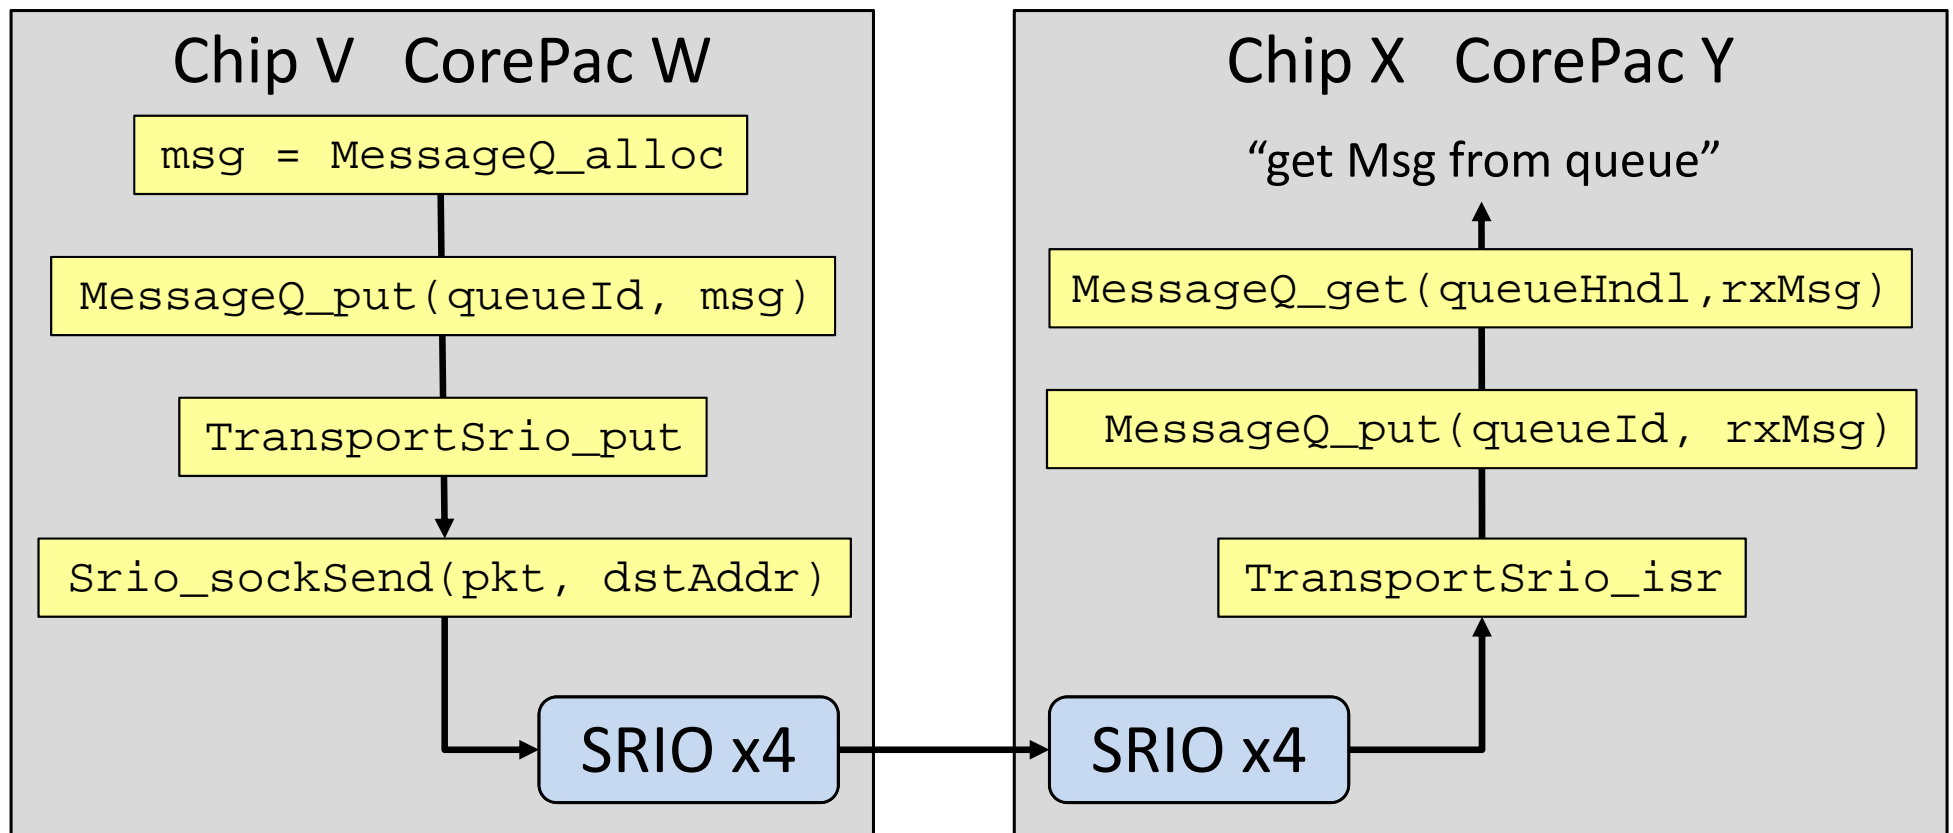
IPC Transports – SRIO (2/3)

- ◆ From a messageQ standpoint, the SRIO transport works the same as the QMSS transport. At the transport level, it is also somewhat the same.
- ◆ The SRIO transport copies the messageQ message into the SRIO data buffer.
- ◆ It will then pop a SRIO descriptor and put a pointer to the SRIO data buffer into the descriptor.



IPC Transports – SRIO (3/3)

- ◆ The transport then passes the descriptor to the SRIO LLD via the `Srio_sockSend` API.
- ◆ SRIO then sends and receives the buffer via the SRIO PKTDMA.
- ◆ The message is then queued on the Receiver side.



Configure the Transport Layer

- ◆ Most of the transport changes are in the CFG file.
- ◆ The XDC tools build the configuration and initialization code. User involvement is only in changing the CFG file.
- ◆ Some additional include and initialization is needed in the code.

Now consider the differences between shared memory and Navigator transport from the application perspective ...

Configuration: Shared Memory CFG File

```
var MessageQ          = xdc.useModule('ti.sdo.ipc.MessageQ');
var Notify             = xdc.module('ti.sdo.ipc.Notify');
var Ipc               = xdc.useModule('ti.sdo.ipc.Ipc');
Notify.SetupProxy      =
    xdc.module(Settings.getNotifySetupDelegate());
MessageQ.SetupTransportProxy=
    xdc.module(Settings.getMessageQSetupDelegate());
```

```
/* Use shared memory IPC */
MessageQ.SetupTransportProxy =
    xdc.module('ti.sdo.ipc.transports.TransportShmSetup');

Program.global.TRANSPORTSETUP =
    MessageQ.SetupTransportProxy.delegate$. $name;
```


Configuration: Navigator CFG File

```
var MessageQ = xdc.useModule('ti.sdo.ipc.MessageQ');  
var Ipc      = xdc.useModule('ti.sdo.ipc.Ipc');  
var TransportQmss = xdc.useModule('ti.transport.ipc.qmss.transports.TransportQmss');
```

```
/* use IPC over QMSS */  
MessageQ.SetupTransportProxy = xdc.useModule(Settings.getMessageQSetupDelegate());  
var TransportQmssSetup =  
    xdc.useModule('ti.transport.ipc.qmss.transports.TransportQmssSetup');  
MessageQ.SetupTransportProxy = TransportQmssSetup;  
  
TransportQmssSetup.descMemRegion = 0;  
Program.global.descriptorMemRegion = TransportQmssSetup.descMemRegion;  
  
Program.global.numDescriptors = 8192;  
  
Program.global.descriptorSize = cacheLineSize; // multiple of cache line size  
  
TransportQmss.numDescriptors = Program.global.numDescriptors;  
  
TransportQmss.descriptorIsInSharedMem = true;  
TransportQmss.descriptorSize = Program.global.descriptorSize;
```

Configuration: SRIO CFG File

```
var MessageQ = xdc.useModule('ti.sdo.ipc.MessageQ');  
var Ipc = xdc.useModule('ti.sdo.ipc.Ipc');  
var TransportSrio = xdc.useModule('ti.transport.ipc.srio.transports.TransportSrio');
```

```
/* use IPC over SRIO */  
MessageQ.SetupTransportProxy = xdc.useModule(Settings.getMessageQSetupDelegate());  
var TransportSrioSetup =  
    xdc.useModule('ti.transport.ipc.srio.transports.TransportSrioSetup');  
MessageQ.SetupTransportProxy = TransportSrioSetup;  
  
TransportSrioSetup.messageQHeapId = 1; /* Sized specifically to handle receive side packets.  
                                         * Heap should  
                                         * not be used by any other application or module */  
TransportSrioSetup.descMemRegion = 0;  
TransportSrioSetup.numRxDescBufs = 256; /* Should be sized large enough so that multiple  
                                         * packets can be queued on receive side and still  
                                         * have buffs available for incoming packets */
```

Configuration: Navigator Initialization (1/2)

Additional Include:

```
/* QMSS LLD*/  
#include <ti/drv/qmss/qmss_drv.h>  
#include <ti/drv/qmss/qmss_firmware.h>  
  
/* CPPI LLD */  
#include <ti/drv/cppi/cppi_drv.h>  
  
#include <ti/transport/ipc/examples/common/bench_common.h>  
  
#include <ti/transport/ipc/qmss/transport/TransportQmss.h>
```

Configuration: Navigator Initialization (2/2)

Add an initialization routine:

```
Int main(Int argc, Char* argv[])
{

    if (selfId == 0)
    {
        /* QMSS, and CPPI system wide initializations are run on
         * this core */
        result = systemInit();
        if (result != 0)
        {
            System_printf("Error (%d) while initializing QMSS\n", result);
        }
    }

    // Note that systemInit is provided by TI
```

IPC Transport Details

Throughput (Mb/second)

Message Size (Bytes)	Shared Memory	Multicore Navigator	SRIO	
48	23.8	34.6	4.1	
256	125.8	184.4	21.2	
1024	503.2	737.4	-	

Benchmark Details

- IPC Benchmark Examples from MCSDK
- CPU Clock – 1 GHz
- Header Size– 32 bytes
- SRIO – Loopback Mode
- Messages allocated up front

IPC Transport Pros/Cons

	PROS	CONS
Shared Memory	<ul style="list-style-type: none">• Simplest to implement• Moderate throughput	<ul style="list-style-type: none">• Single device only• Requires Notify module and API call if doorbell required• Possible contention with other tasks using the same shared memory
Multicore Navigator	<ul style="list-style-type: none">• Highest throughput• Dedicated resources• Consumes least CPU cycles• Interrupt generated when data is available	<ul style="list-style-type: none">• Single device only
SRIO	<ul style="list-style-type: none">• Can be used across devices	<ul style="list-style-type: none">• Lowest throughput

For More Information

- For more information, refer to the [BIOS MCSDK User Guide](#) on the [Embedded Processors Wiki](#).
- For questions regarding topics covered in this training, visit the support forums at the [TI E2E Community](#) website.